

Python For Data Mining



PYTHON PACKAGES FOR DATA MINING

NUMPY

NumPy is the fundamental package for scientific computing with Python. NumPy is an extension to the Python programming language, adding support for large, multi-dimensional arrays and matrices, along with a large library of high-level mathematical functions to operate on these arrays

SCIPY

SciPy (pronounced “Sigh Pie”) is open-source software for mathematics, science, and engineering. The SciPy library is built to work with NumPy arrays, and provides many user-friendly and efficient numerical routines such as routines for numerical integration and optimization.

Pandas

Pandas is a Python package providing fast, flexible, and expressive data structures designed to make working with “relational” or “labeled” data both easy and intuitive. It aims to be the fundamental high-level building block for doing practical, real world data analysis in Python.

MATPLOTLIB

matplotlib is a plotting library for the Python programming language and its NumPy numerical mathematics extension. It provides an object-oriented API for embedding plots into applications using general-purpose GUI toolkits.

Scalars, Vectors and Matrices

Scalar

24

Vector

$\begin{bmatrix} 2 & -8 & 7 \end{bmatrix}$

row

or
column

$\begin{bmatrix} -6 \\ -4 \\ 27 \end{bmatrix}$

Matrix

$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}$

row(s) \times column(s)

conda install -c anaconda numpy

Numpy

Numpy is the core library for scientific computing in Python. It provides a high-performance multidimensional array object, and tools for working with these arrays.

```
import numpy as np  
a = np.array([1, 2, 3])  
print(type(a))  
print(a.shape)  
print(a[0], a[1], a[2])  
a[0] = 5  
print(a)
```

```
b = np.array([[1,2,3],[4,5,6]])  
print(b.shape)  
print(b[0, 0], b[0, 1], b[1, 0])
```

```
import numpy as np  
a=np.array([1,2,3,4,5,6],float)  
print(a)  
print(type(a))
```

```
import numpy as np
a=np.array([1,2,3,4,5,6],float)
print(a)
print(type(a))
print(a[3])
print(a[:3])
a[3]=54
print(a)
```

Array

```
import numpy as np  
array1=np.array([2,3,4,5,6])  
array2=array1*2  
print(array1)  
print(array2)
```

```
import numpy as np
a=np.array([1,2,3,4,5,6,7,8,9],float)
print(a)
print(a.shape)
print(all(a))
print(any(a))
print(a[:1])
print(a[:3])
print(a[1:])
print(a[3:])
print(a[:-1])
```



```
import numpy as np
a=np.array([[1,2,3,4,5,6],[9,11,22,33,44,55]],float)
print(a)
print(type(a))
print(a[0,1])
print(a[1,1])
print(a[1:,])
print(a[:,2])
print(a.shape)
print(a.dtype)
print(len(a))
print(34 in a)
print(22 in a)
```

```
import numpy as np
a=np.array([2,3,4],float)
print(a)
s=a.tostring()
print(s)
print(np.fromstring(s))
a.fill(0)
print(a)
```

```
import numpy as np
a = np.array(range(6), float).reshape((2, 3))
print(a)
a.transpose()
print(a)
```

```
import numpy as np
a=np.array([1,2],float)
b=np.array([3,4,5],float)
c=np.array([6,7,8,9,0])
d=np.concatenate((a,b,c),float)
print(a)
print(b)
print(c)
print(d)
```

```
import numpy as np
a=np.array([[1,2],[3,4]],float)
b=np.array([[5,6],[7,8]],float)
c=np.concatenate((a,b))
print(c)
d=np.concatenate((a,b),axis=1)
print(d)
```

Numpy functions to create arrays

```
import numpy as np  
print(np.eye(4))
```

```
import numpy as np  
d=np.ones(4)  
print(type(d))
```

```
import numpy as np  
d=np.zeros(3)  
d
```

```
c = np.full((2,2), 7)  
print(c)
```

Basic array operations

```
import numpy as np
a=np.array([1,2,4,5,68,9,0],float)
print(a)
print(a.sum())
print(a.prod())
print(a)
print(a.mean())
print(a.var())
print(a.std())
print(a.max())
print(a.min())
print(a.argmax())
print(a.argmin)
```

```
print(np.sum(a))
```

Arange

With NumPy we can get an array based on ranges. The arange method receives 1, 2 or 3 arguments

```
import numpy as np
array1=np.arange(5)
print(array1)
array1=np.arange(5,10)
print(array1)
array1=np.arange(0,10,2)
print(array1)
```

Random numbers

```
import numpy as np
a=np.random.rand(5)
print(a)
b=np.random.rand(2,3)
print(b)
c=np.random.random()
print(c)
d=np.random.randint(2,45)
print(d)
```

ndarray data type

Type	Description
bool	Boolean (True or False) stored as a bit
int8	Byte (-128 to 127)
int16	Integer (-32768 to 32767)
int32	Integer (-2^{31} to $2^{31} - 1$)
int64	Integer (-2^{63} to $2^{63} - 1$)
uint8	Unsigned integer (0 to 255)
uint16	Unsigned integer (0 to 65535)
uint32	Unsigned integer (0 to $2^{32} - 1$)
uint64	Unsigned integer (0 to $2^{64} - 1$)
float16	Half precision float: sign bit, 5b expo, 10b mantissa
float32	Single precision float: sign bit, 8b expo, 23b mantissa
float64	Double precision float: sign bit, 11b expo, 52b mantissa
complex64	Complex number, represented by two 32-bit floats (real & imag)
complex128	Complex number, represented by two 64-bit floats (real & imag)

Popular methods for working with ndarrays

Function	Description	Example
Logical		
<code>all()</code>	True if all elements are nonzero	<pre>>>> x = array([0,1,2,0,4,5]) all(x) = False</pre>
<code>any()</code>	True if any (at least one) elements are nonzero	<pre>>>> x = array([0,1,2,0,4,5]) any(x) = True</pre>
<code>find()</code>	Return the indices where <i>ravel</i> (condition) is true	<pre>>>> x = array([0,1,2,2,1,7]) find(x >= 3) = array([5])</pre>
Slicing 1D arrays (a few cases)		
<code>x[n:m]</code> <code>x[:m]</code> <code>x[n:]</code> <code>x[n:-1]</code> <code>x[n:-2]</code> <code>x[n:m:k]</code>	The 1D subarray from n to m-1 The 1D subarray from 0 to m-1 The 1D subarray from n to the end The 1D subarray from n to end-1 The 1D subarray from n to end-2 The 1D subarray from n to m-1 with k index striding	<pre>>>> x = array([0,1,2,3,4,5]) x[:2] = array([0,1]) x[:3] = array([0,3]) x[1:-2] = array([1,2])</pre>
Slicing 2D arrays (a few cases)		
<code>x[n:m,j:k]</code> <code>x[n:m,:]</code> <code>x[:,j:k]</code> <code>x[n:m:o,j:k:l]</code> <code>x[n:-1,:]</code> <code>x[:,j:-2]</code> <code>x[3,:]</code> <code>x[:,0]</code>	The 2D subarray from n to m-1, j to k-1 The 2D subarray from 0 to m-1, all columns The 2D subarray all rows, columns j to k-1 The 2D subarray with striding by o and l in rows and columns respectively The 2D subarray from n to end-1, all columns The 2D subarray all rows, columns j to k-2 The 2D subarray row 3, all columns The 2D subarray all rows, column 0	<pre>>>> x = array([[0,1,2], [3,4,5]]) x[:2,:2] = array([[0,1],[3,4]]) x[-1,-1] = array([[5]])</pre>

Function	Description	Example
Shape & Concatenation		
<code>reshape()</code>	Reshape a 1D or 2D array to a new shape; the new shape must be consistent.	<pre>>>> x = arange(0,5) 1D 6 elements y = reshape(x,(2,3)) 2D 2x3 elements</pre>
<code>concatenate()</code>	Join a sequence of arrays together. The arrays must have the same shape except in the axis used for combining. axis=0 is rows, axis=1 is columns.	<pre>>>> x = array([[0,1,2,3,4,5]]) 2D 1x6 elements concatenate((x,x), axis=0) 2D 1x6 elements concatenate((x,x)), axis=1) 2D 1x12 elements</pre>
<code>hstack()</code>	Stack arrays horizontally. A subset of <code>concatenate()</code>	<pre>>>> x = array([[0,1,2,3,4,5]]) 2D 1x6 elements x = x.T #transpose y=hstack((x,x)) 2D 6x2 columns</pre>
<code>vstack()</code>	Stack arrays vertically. A subset of <code>concatenate()</code>	<pre>>>> x = array([[0,1,2,3,4,5]]) 2D 1x6 elements y=vstack((x,x)) 2D 2x6 columns</pre>
<code>flatten()</code>	Values of the argument array become a 1D array. May be done in-place with <code>x.flatten()</code>	<pre>>>> x = array([[0,1,2,3,4,5]]) x.flatten() 1D 6 element</pre>
<code>transpose()</code> or <code>array.T</code>	Like matrix transpose for 2D arrays. In-place via <code>x.T</code> .	<pre>>>> x = array([[0,1,2,3,4,5]]) x.T 2D 6x1 array</pre>

Math	Many other standard functions, e.g., trig, are also available for array operations	
mean(x)	The sample mean of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) mean(x) = 2.5
var(x)	The sample variance of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) var(x) = 2.9167
std(x)	The sample standard deviation of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) std(x) = 1.7078
sum(x)	The sum of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) sum(x) = 15
prod(x)	The sample mean of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) prod(x) = 0
cumsum(x)	The sample mean of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) cumsum(x) = array([0, 1,3,6,10,15])
cumprod(x)	The sample mean of the values contained in array x.	>>> x = array([1,1,2,3,4,5]) cumprod(x) = array([1, 1,2,6,24,120])
min(x)	The sample mean of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) min(x) = 0
max(x)	The sample mean of the values contained in array x.	>>> x = array([0,1,2,3,4,5]) max(x) = 5
conj(x)	The sample mean of the values contained in array x.	>>> x = array([2+5j]) conj(x) = array([2+5j])
x.real & x.imag	The real or imaginary part of the values contained in array x. Also real(x), imag(x)	>>> x = array([2+5j]) x.real = array([2.]) imag(x) = array([5.])

$$10 \begin{pmatrix} 3 & 3 \end{pmatrix} \begin{pmatrix} 2 \\ 2 \end{pmatrix}$$

```
import numpy as np  
m1=np.array([3.,3.])  
m2=np.array([2.,2.])  
print(10 * np.dot(m1,m2))
```

Matplotlib

The library is generally used as follows:

1. Call a plotting function with some data (e.g. `plot()`).
2. Call many functions to setup the properties of the plot (e.g. labels and colors).
3. Make the plot visible (e.g. `show()`).

```
from matplotlib import pyplot as plt  
plt.plot([1,2,3],[4,5,1])  
plt.show()
```

```
import matplotlib.pyplot as plt plt.plot([1,2,3,4], [1,4,9,16],  
'ro')  
plt.axis([0, 6, 0, 20]) plt.show()
```

```
from matplotlib import pyplot as plt from matplotlib  
import style  
import numpy as np  
style.use('ggplot')  
x,y = np.loadtxt('exampleFile.csv', unpack=True,  
delimiter = ',')  
plt.plot(x,y)  
plt.title('Epic Info')  
plt.ylabel('Y axis')  
plt.xlabel('X axis')  
plt.show()
```

```
import matplotlib.pyplot as plt
import numpy
myarray = numpy.array([1, 2, 3])
plt.plot(myarray)
plt.xlabel('some x axis')
plt.ylabel('some y axis')
plt.show()
```

```
import matplotlib.pyplot as plt
import numpy
x = numpy.array([1, 2, 3])
y = numpy.array([2, 4, 6])
plt.scatter(x,y)
plt.xlabel('some x axis')
plt.ylabel('some y axis')
plt.show()
```

Pandas data types

Series

DataFrame

Panel

Panel4D

```
import pandas as pd  
c=pd.Series([1,2,3,4,5])  
c
```

```
c.values
```


Creating DataFrame

	a	b	c
1	4	7	10
2	5	8	11
3	6	9	12

```
import pandas as pd
df=pd.DataFrame({"a":[4,5,6],"b":[7,8,9],"c":[10,11,12]},
index=[1,2,3])
```

```
df=pd.DataFrame([[4,7,10],[5,8,11],[6,9,12]],index=[1,2,3],
columns=["a","b","c"])
```

```
states = {'State' : ['Gujarat', 'Tamil Nadu', ' Andhra', 'Karnataka',  
'Kerala'], 'Population': [36, 44, 67,89,34], 'Language' : ['Gujarati',  
'Tamil', 'Telugu', 'Kannada', 'Malayalam']}
```

```
State=pd.DataFrame(states)  
State
```

```
df=pd.DataFrame({"x1":["a','b','c'], "x2":[1,2,3]})  
df
```

```
df1=pd.DataFrame({"x1":["A','B','C'], "x3":["T","F","T"]})  
df1
```

```
pd.merge(df,df1,how='left',on='x1')
```

```
pd.merge(df,df1,how='right',on='x1')
```

```
pd.merge(df,df1,how='inner',on='x1')
```

```
pd.merge(df,df1,how='outer',on='x1')
```

Pandas indexing

For indexing the rows loc, iloc and ix

A.loc[3,:]

A.ix[2,:]

A.iloc[2,:]

A.loc[:, 'b']

	Index	a	b	c
A =	2	4	5	yes
	3	6.5	7	no
	6	8	9	ok

```
import pandas as pd
```

```
A = pd.DataFrame([[4, 5, 'yes'], [6.5, 7, 'no'], [8, 9, 'ok']],  
index=[2, 3, 6], columns=['a', 'b', 'c'])
```

```
import pandas as pd
data = pd.DataFrame({'group': ['a', 'a', 'a', 'b','b', 'b', 'c', 'c','c'],
                     'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
data
```

```
data.describe()
```

df.info() shows data types, number of rows and columns, and memory usage of your data frame

Sorting

```
data.sort_values(by=['group','ounces'], ascending=[False,
True], inplace=True)
data
```

Removing duplicates

```
import pandas as pd  
data = pd.DataFrame({'k1': ['one'] * 3 + ['two'] * 4, 'k2': [3, 2, 1, 3,  
3, 4, 4]}))
```

```
data.drop_duplicates()  
data.drop_duplicates(subset='k1')
```

Creating a new column based on values from another column

```
import pandas as pd  
data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',  
'Pastrami','corned beef', 'Bacon', 'pastrami', 'honey ham','nova  
lox'],'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]}))
```

	food	ounces	animal
0	bacon	4.0	pig
1	pulled pork	3.0	pig
2	bacon	12.0	pig
3	Pastrami	6.0	cow
4	corned beef	7.5	cow
5	Bacon	8.0	pig
6	pastrami	3.0	cow
7	honey ham	5.0	pig
8	nova lox	6.0	salmon

```
meat_to_animal = {  
    'bacon': 'pig',  
    'pulled pork': 'pig',  
    'pastrami': 'cow',  
    'corned beef': 'cow',  
    'honey ham': 'pig',  
    'nova lox': 'salmon'  
}
```

```
def meat2animal(series):  
    if series["food"]=='bacon':  
        return 'pig'  
    elif series["food"]=='pulled pork':  
        return 'pig'  
    elif series["food"]=='pastrami':  
        return 'cow'  
    elif series["food"]=='corned beef':  
        return 'cow'  
    elif series["food"]=='honey ham':  
        return 'pig'  
    else:  
        return 'salmon'
```

```
data['animal'] =  
data['food'].map(str.lower).map(meat_to_animal)  
data
```


	food	ounces	animal	animal2
0	bacon	4.0	pig	pig
1	pulled pork	3.0	pig	pig
2	bacon	12.0	pig	pig
3	Pastrami	6.0	cow	salmon
4	corned beef	7.5	cow	cow
5	Bacon	8.0	pig	salmon
6	pastrami	3.0	cow	cow
7	honey ham	5.0	pig	pig
8	nova lox	6.0	salmon	salmon

```
data['animal2'] = data.apply(meat2animal,axis='columns')
data
```

Removing or dropping a column

```
data.drop('animal', axis='columns', inplace=True)  
data
```

```
import pandas as pd  
import numpy as np  
data = pd.Series([1., -999., 2., -999., -1000., 3.])  
data
```

```
data.replace(-999, np.nan, inplace=True)  
data
```

```
import pandas as pd  
data = pd.Series([1., -999., 2., -999., -1000., 3.])  
data
```

```
data.replace([-999, -1000], np.nan, inplace=True)  
data
```

Dropping NaN

```
data.dropna()
```

Filling/replacing NaN values with something else (replace NaN with 0/zero)

```
data.dropna()
```

Create dataframe

	first_name	last_name	age	preTestScore	postTestScore
0	Jason	Miller	42	4	25,000
1	Molly	Jacobson	52	24	94,000
2	Tina	.	36	31	57
3	Jake	Milner	24	.	62
4	Amy	Cooze	73	.	70

```
raw_data = {'first_name': ['Jason', 'Molly', 'Tina', 'Jake', 'Amy'],  
            'last_name': ['Miller', 'Jacobson', ".", 'Milner', 'Cooze'],  
            'age': [42, 52, 36, 24, 73],  
            'preTestScore': [4, 24, 31, ".", "."],  
            'postTestScore': ["25,000", "94,000", 57, 62, 70]}
```

```
df = pd.DataFrame(raw_data, columns = ['first_name',  
                                       'last_name', 'age', 'preTestScore', 'postTestScore'])  
df
```

Save dataframe as csv in the working director

```
df.to_csv('C:/Users/shubh/Desktop/python/data.csv')
```

Load a csv

```
df=pd.read_csv('C:/Users/shubh/Desktop/python/data.csv')
```

Load a csv with no headers

```
df=pd.read_csv('C:/Users/shubh/Desktop/python/data.csv')  
df
```

Pivot table

In data processing, a ***pivot table*** is a data summarization tool found in data visualization programs such as spreadsheets or business intelligence software. Among other functions, a pivot table can automatically sort, count, total or average the data stored in one table or spreadsheet, displaying the results in a second table showing the summarized data.

```
pandas.pivot_table(data, values=None, index=None,  
columns=None, aggfunc='mean', fill_value=None, margins=False,  
dropna=True, margins_name='All')
```

EXCEL/CSV File

```
import pandas as pd  
df=pd.read_excel('C:/Users/shubh/Desktop/data/sales-funnel.xlsx')
```

```
import pandas as pd  
df=pd.read_csv('C:/Users/shubh/Desktop/data/sales-funnel.xlsx')
```

```
df.head()
```



```
import pandas as pd
import numpy as np
df=pd.read_excel('C:/Users/shubh/Desktop/data/
excel-comp-data.xlsx')
df.head()
```

	account	name	street	city	state	postal- code	Jan	Feb	Mar
0	211829	Kerluke, Koepp and Hilpert	34456 Sean Highway	New Jaycob	Texas	28752	10000	62000	35000
1	320563	Walter-Trantow	1311 Alvis Tunnel	Port Khadijah	NorthCarolina	38365	95000	45000	35000
2	648336	Bashirian, Kunde and Price	62184 Schamberger Underpass Apt. 231	New Lilianland	Iowa	76517	91000	120000	35000
3	109996	D'Amore, Gleichner and Bode	155 Fadel Crescent Apt. 144	Hyattburgh	Maine	46021	45000	120000	10000
4	121213	Bauch-Goldner	7274 Marissa Common	Shanahanchester	California	49681	162000	120000	35000

```
df["total"] = df["Jan"] + df["Feb"] + df["Mar"]  
df.head()
```

Performing column level analysis is easy in pandas

```
df["Jan"].sum(), df["Jan"].mean(),df["Jan"].min(),df["Jan"].max()
```

```
sum_row=df[["Jan","Feb","Mar","total"]].sum()  
sum_row
```

```
df_sum=pd.DataFrame(data=sum_row).T  
df_sum
```

PANDAS Example

```
import numpy as np
import matplotlib as p
import pdb
from pandas import *
data = read_csv('C:/Users/shubh/Desktop/data/train.csv')
data
```

```
data.describe()
```

```
import numpy as np
import matplotlib as p
import pdb
from pandas import *
data = read_csv('C:/Users/shubh/Desktop/data/train.csv')
data
```

```
men = data[data.Sex == 'male']
women = data[data.Sex == 'female']
```

We can then find the proportion of men and women that survive

```
proportion_women_survived =
float(sum(women.Survived))/len(women)
proportion_men_survived = float(sum(men.Survived))/len(men)
```

```
data.groupby('Sex').Survived.mean()
```

```
men.Age  
women.Age
```

```
data.columns
```

Now we might want to create a new column that stores our prediction for whether someone survived or not. Let's call this new column 'prediction'

```
data['prediction']=0
```

```
data.prediction[data.sex == 'female']=1
```

Visualization

```
import matplotlib.pyplot as plt  
plt.plot([1,2,3,4])  
plt.ylabel('some numbers')  
plt.show()
```

```
plt.plot([1, 2, 3, 4], [1, 4, 9, 16])  
plt.show()
```

```
import matplotlib.pyplot as plt  
plt.plot([1,2,3,4], [1,4,9,16], 'ro')  
plt.axis([0, 6, 0, 20])  
plt.show()
```

```
#import matplotlib library
import matplotlib.pyplot as plt
#define some data
x = [1,2,3,4]
y = [20, 21, 20.5, 20.8]
#plot data
plt.plot(x, y)
#show plot
plt.show()
```


Regression analysis using Python

Linear regression analysis means “fitting a straight line to data”

- also called linear modelling

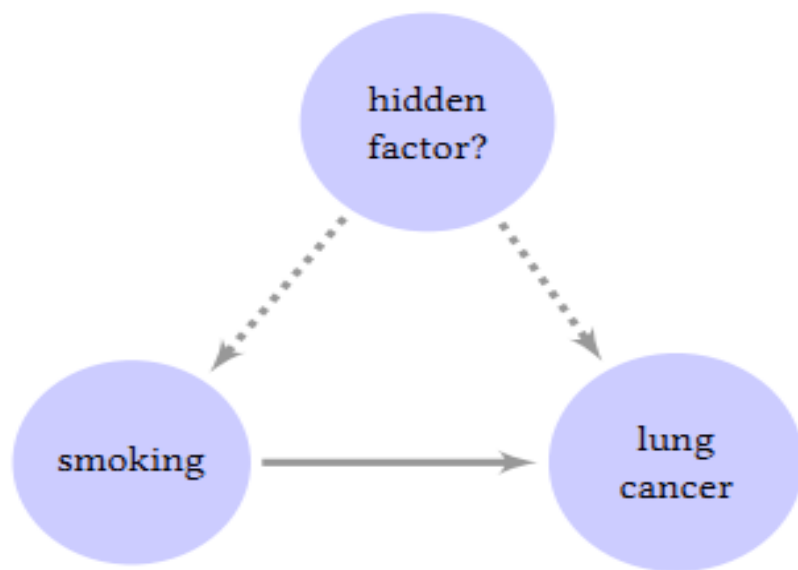
It's a widely used technique to help model and understand real-world phenomena

- easy to use
- easy to understand intuitively

Allows prediction

cigarettes smoked (per day)	CVD mortality (per 100 000 men per year)	lung cancer mortality (per 100 000 men per year)
0	572	14
10 (actually 1-14)	802	105
20 (actually 15-24)	892	208
30 (actually >24)	1025	355

```
import pandas
import matplotlib.pyplot as plt
data = pandas.DataFrame({'cigarettes': [0,10,20,30],
'CVD': [572,802,892,1025],
'lung': [14,105,208,355]});
data.plot('cigarettes', 'CVD', kind='scatter')
plt.title("Deaths for different smoking intensities")
plt.xlabel("Cigarettes smoked per day")
plt.ylabel("CVD deaths")
plt.show()
```



```
import numpy, pandas
import matplotlib.pyplot as plt
import statsmodels.formula.api as smf
df = pandas.DataFrame({'cigarettes': [0,10,20,30],
'CVD': [572,802,892,1025],
'lung': [14,105,208,355]});
df.plot('cigarettes', 'CVD', kind='scatter')
lm = smf.ols("CVD ~ cigarettes", data=df).fit()
xmin = df.cigarettes.min()
xmax = df.cigarettes.max()
X = numpy.linspace(xmin, xmax, 100)
# params[0] is the intercept ( $\beta_0$ )
# params[1] is the slope ( $\beta_1$ )
Y = lm.params[0] + lm.params[1] * X
plt.plot(X, Y, color="darkgreen")
plt.show()
```

```
ipython qtconsole --pylab=inline
```